


## P4c – Deep CNN Models

**Jorge Henriques**

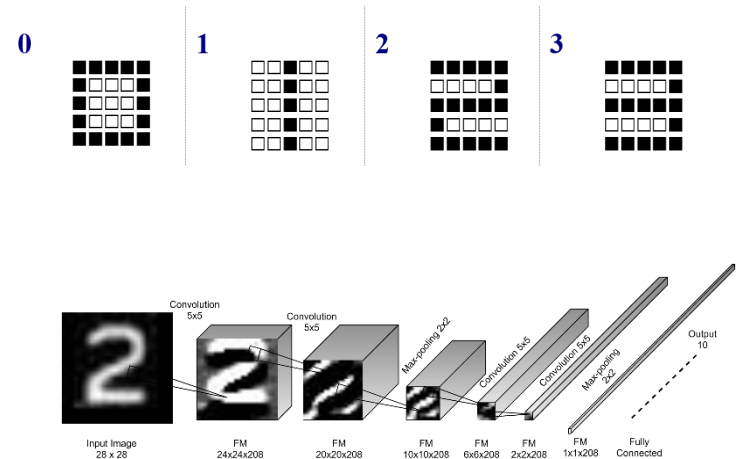
jh@dei.uc.pt

Departamento de Engenharia Informática  
Faculdade de Ciências e Tecnologia



UNIVERSIDADE D  
COIMBRA

**dei** engenharia  
informática



# Contents

- 1 | Objectives
- 2 | Datasets
- 3 | Tasks
- 4 | Conclusions

## ■ Deep learning Concepts

- 1| Implement Digit recognition from scratch
  - One layer network
  - Backpropagation
  
- 2| Use python/keras learning functionalities
  - Deep neural network
    - MLNN / autoencoders
  - Convolutional neural network
  
- 3| Evaluation the performance of the DNN/CNN classifier


# Contents

- 1 | Objectives
- 2 | Datasets
- 3 | Tasks
- 4 | Conclusions

## ■ Datasets

### • 1 | JH dataset

- `digitsX.csv`
- `digitsT.csv`
- Digits (5,5)
- N=100 (10 examples of each digit)



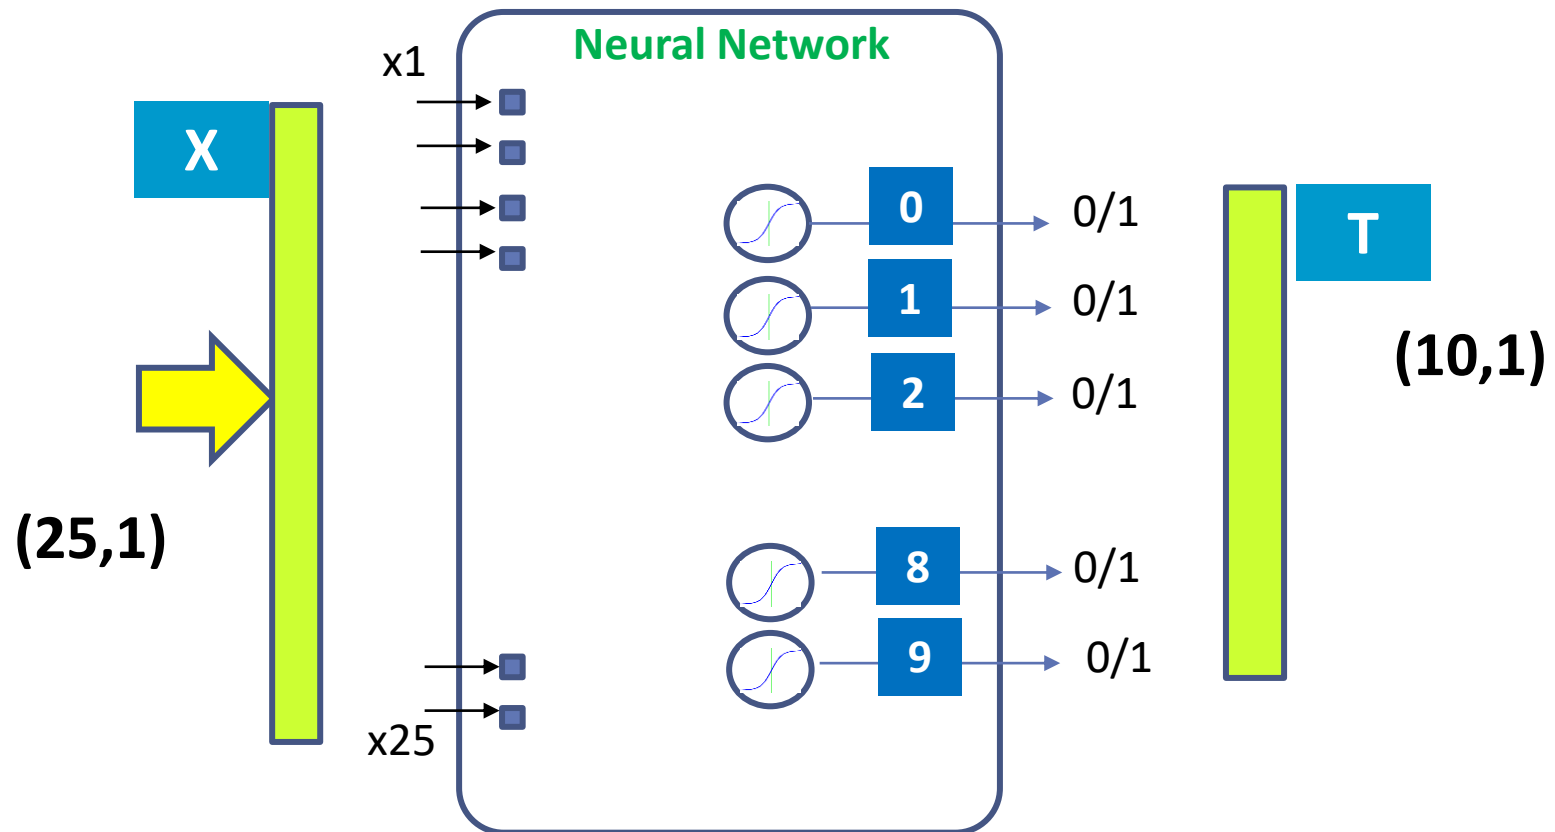
### • 2 | MNIST dataset

- `from keras.datasets import mnist`
- `X = mnist.load_data()`
- Digits (28,28)
- N=70000 (60000 train+10000 test)



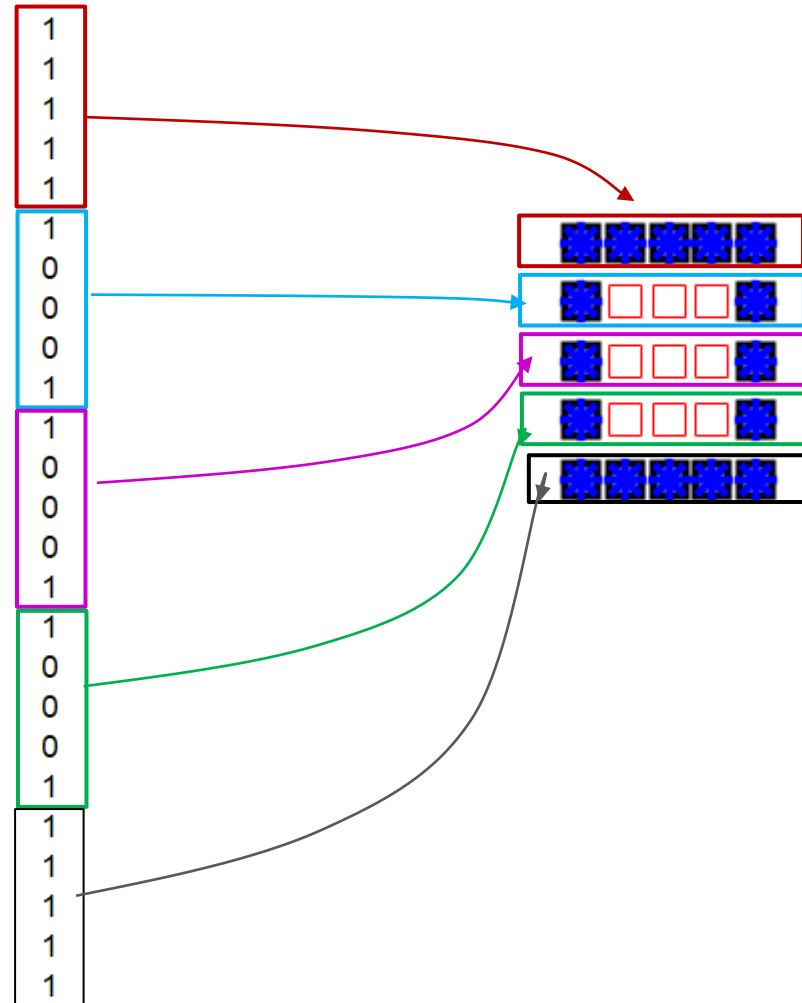

## ■ Neural network : structure ?

- Inputs (**vector**) / outputs (**one node per class**)



## Inputs

- Image (matrix) → vector
- (5x5) → (25,1)
- $X =[:,1]$
- (25,1)



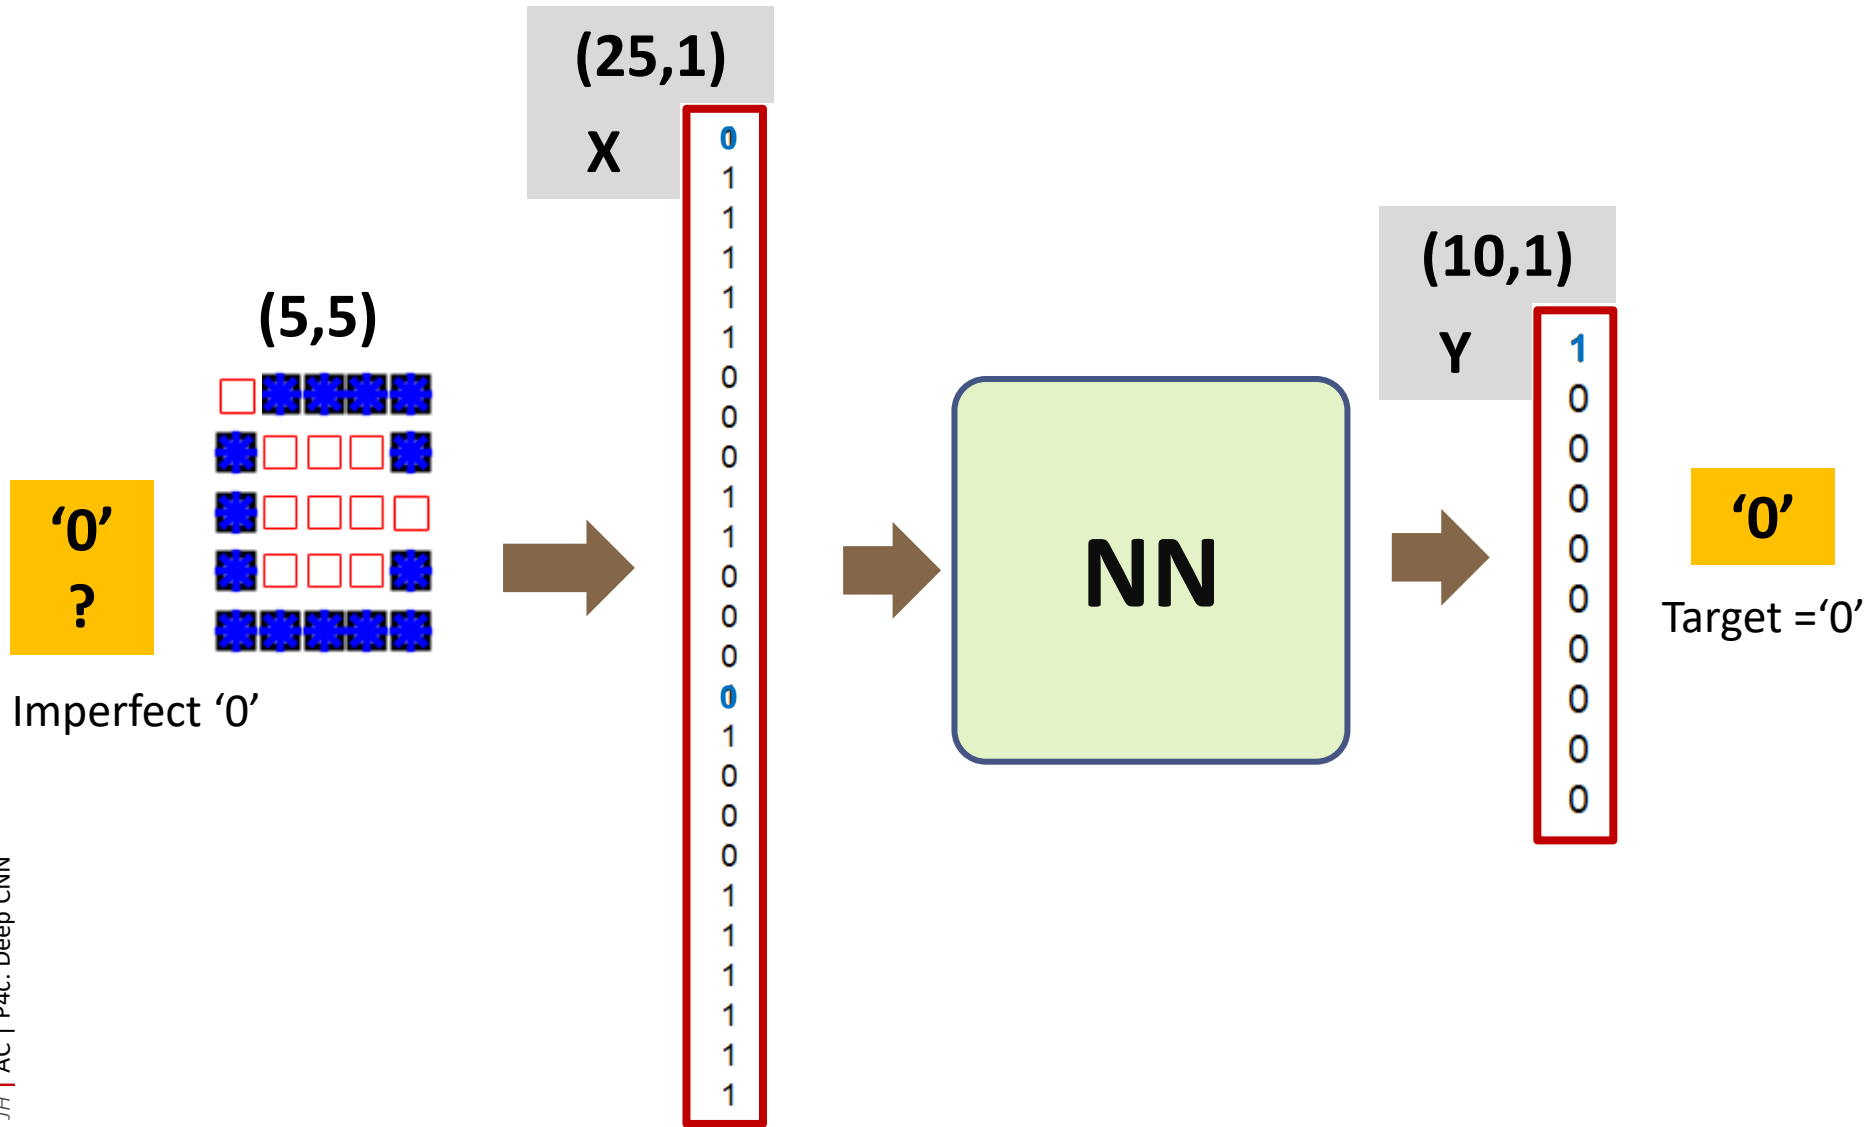
## ■ Target

- One node per class: (10,1)

Digit	0	1	2	3	4	5	6	7	8	9
Output	1	0	0	0	0	0	0	0	0	0
	0	1	0	0	0	0	0	0	0	0
	0	0	1	0	0	0	0	0	0	0
	0	0	0	1	0	0	0	0	0	0
	0	0	0	0	1	0	0	0	0	0
	0	0	0	0	0	1	0	0	0	0
	0	0	0	0	0	0	1	0	0	0
	0	0	0	0	0	0	0	1	0	0
	0	0	0	0	0	0	0	0	1	0
	0	0	0	0	0	0	0	0	0	1



### ■ One example





### ■ MNIST dataset


- Total of 70,000 examples of handwritten digits
- Divided into two subsets:
  - Training set: 60,000 examples
  - Test set: 10,000 examples
- Each image is grayscale [0 .. 255] image of size 28×28 pixels (784 total features for flat/vector input).
- Labeled with one of 10 classes (digits 0.. 9).

# Contents

- 1 | Objectives
- 2 | Datasets
- 3 | Tasks
- 4 | Conclusions

## ■ JH dataset

- Adaline ?
- One hidden layer
- MLNN (one hidden layer) ?



# Last Class


- SE, SP, F1score

## Computed for each digit


- For example, with respect to digit '2'
- $T = [2 \ 2 \ 7 \ 3 \ 4 \ 5 \ 6 \ 7 \ 0 \ 9 \ 9 \ 2 \ 0]$
- $T1 = [1 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0]$
- $Y_{net} = [0 \ 1 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 1]$
-

## ■ Task 1: One-layer or MLNN (hidden layer)

```
#----- DATA X(25,100) T(10,100)
filename = 'digitsX.csv'
df       = read_csv(filename)
X        = df.values
filename = 'digitsT.csv'
df       = read_csv(filename)
T        = df.values
N        = X.shape[1]
```



```
#----- Image visualization
dig = X[:,36]          # dimension (25,1)
plt.figure(1)
plt.imshow( dig.reshape( (5,5)), cmap=plt.cm.gray_r)
plt.show()
```



- **Task 2: Deep NN**
  - MNIST dataset




- Total of 70,000 examples of handwritten digits
- Divided into two subsets:
  - Training set: 60,000 examples
  - Test set: 10,000 examples
- Each image is grayscale image of size 28×28 pixels (784 total features for flat input).
- Labeled with one of 10 classes (digits 0–9).

- **Task 2: Implement a Deep NN**




- 2.1 | MLNN – “high” number of layers



- 2.2 | Encoders - decoders

- 2.3 | CNN - Convolutional neural networks

## ■ Task 2.1 | Deep NN - MLNN



```
from keras.datasets import mnist
#----- load MNIST [0-255] gray scale
#---- X = mnist.load_data()
```

```
(Xtrain, Ttrain), (Xtest, Ttest) = mnist.load_data()
```

```
> Xtrain.shape
> (60000, 28, 28)
```


```
#----- Reshape
Xtrain = Xtrain.reshape(Xtrain.shape[0], num_pixels).astype('float32')
Xtest = Xtest.reshape(Xtest.shape[0], num_pixels).astype('float32')
```

```
> Xtrain.shape
> (60000, 784)
```

```
#----- Normalization [0.155] -> [0..1] (gray->bw)
Xtrain = Xtrain / 255
```



## ■ Task 2.1 | Deep NN - MLNN




```
#----- Target [6000,1]
(Xtrain, Ttrain), (Xtest, Ttest) = mnist.load_data()

> Ttrain.shape
> (60000,)          vector
```

```
from tensorflow.keras.utils import to_categorical
#----- Target [6000.10] - output - ten neurons
Ttrain = to_categorical(Ttrain)

> Ttrain.shape
> (60000,10)        matrix
```

## ■ Task 2.1 | Deep NN - MLNN



```
from keras.models import Sequential
from keras.layers import Dense
#----- Define and Train [6000.10]
Numclasses=10

model = Sequential()
model.add( Dense(19, activation='relu') )
model.add( Dense(5, activation='relu') )
model.add( Dense(4, activation='relu') )
model.add( Dense(3, activation='relu') )
model.add( Dense(num_classes, activation='softmax') )

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

model.fit( Xtrain, Ttrain, epochs=12, batch_size=500, verbose=2)

Ytrain = model.predict(Xtrain).round()
Ytest  = model.predict(Xtest).round()
```



- Softmax layer as the output layer
- For *classification* problems it is advantageous to use the activation function *softmax*, at the output layer
- *Softmax* normalizes the output and **mimic a probability** vector for each output.
- Softmax function accounts not only for the weighted sum of the inputs, but also for the inputs to the other output nodes

$$f(x_i) = \frac{e^{x_i}}{\sum_{i=1}^n e^{x_i}}$$

- Softmax layer as the output layer

### *Output Layer*

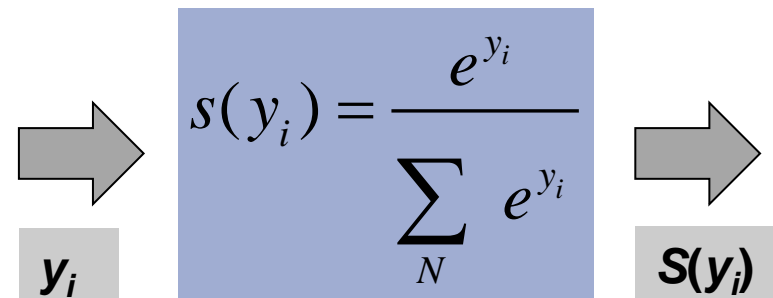
$$h_1 \longrightarrow \sigma \longrightarrow y_1 = \sigma(h_1)$$

$$h_2 \longrightarrow \sigma \longrightarrow y_2 = \sigma(h_2)$$

$$h_3 \longrightarrow \sigma \longrightarrow y_3 = \sigma(h_3)$$

In general, the output of network can be any value.


May not be easy to interpret




The diagram illustrates the Softmax function. On the left, a gray box labeled  $y_i$  has a large gray arrow pointing to a blue box. Inside the blue box is the formula  $s(y_i) = \frac{e^{y_i}}{\sum_N e^{y_i}}$ . Another large gray arrow points from the blue box to a gray box on the right labeled  $S(y_i)$ .

- Softmax layer as the output layer

### Softmax Layer



## ■ Task 2.1: MLNN



- N\_EPOCHS=2

Test	7	1	0	1	5	3	8	6	0	1
Target	1	1	7	1	8	8	9	1	7	1


- N\_EPOCHS=12

Test	3	5	2	9	6	4	9	9	7	4
Target	3	5	2	9	7	4	9	9	7	4


- N\_EPOCHS=32

Test	1	6	1	0	3	5	0	1	7	7
Target	1	6	1	2	3	5	0	1	7	7

## ■ Task 2.2: Autoencoders



## ■ Task 2.2: Autoencoders



```
(Xtrain, Ttrain), (Xtest, Ttest) = mnist.load_data()
```

```
#----- normalization, reshape
```

```
X_train = X_train.astype('float32') / 255.
```


```
X_test = X_test.astype('float32') / 255.
```

```
X_train = np.reshape(X_train, (len(X_train), 28, 28, 1))
```

```
X_test = np.reshape(X_test, (len(X_test), 28, 28, 1))
```



## ■ Task 2.2: Autoencoders



```
#----- ENCODER
input_image = Input(shape =(28, 28, 1))


x = Conv2D(16, (3, 3), activation ='relu', padding ='same')(input_image)
x = MaxPooling2D((2, 2), padding ='same')(x)

x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(x)
x = MaxPooling2D((2, 2), padding ='same')(x)

x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(x)
encoded_layer = MaxPooling2D((2, 2), padding ='same')(x)
```

- Conv2D(16, (3, 3) -----
  - Creates a 2D convolutional layer with **16 filters** (or kernels).
  - Each filter has a kernel size of 3x3 (**mask**).
  - activation='relu':
  - padding='same': Ensures the output has the same spatial dimensions (height and width) as the input by adding zero-padding around the input

## ■ Task 2.2: Autoencoders



#----- ENCODER

```
input_image = Input(shape =(28, 28, 1))
```

```
x = Conv2D(16, (3, 3), activation ='relu', padding ='same')(input_image)
```

```
x = MaxPooling2D((2, 2), padding ='same')(x)
```

```
x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(x)
```

```
x = MaxPooling2D((2, 2), padding ='same')(x)
```

```
x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(x)
```

```
encoded_layer = MaxPooling2D((2, 2), padding ='same')(x)
```

#----- DECODER

```
x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(encoded_layer)
```

```
x = UpSampling2D((2, 2))(x)
```


```
x = Conv2D(8, (3, 3), activation ='relu', padding ='same')(x)
```

```
x = UpSampling2D((2, 2))(x)
```


```
x = Conv2D(16, (3, 3), activation ='relu')(x)
```

```
x = UpSampling2D((2, 2))(x)
```

```
decoded_layer = Conv2D(1, (3, 3), activation ='sigmoid', padding ='same')(x)
```



## ■ Task 2.2: Autoencoders



```
#----- BUILT
autoencoder = Model(input_image, decoded_layer)

autoencoder.compile(optimizer = 'adam', loss = 'binary_crossentropy')


#----- TRAIN
autoencoder.fit(X_train, X_train, epochs = 2, batch_size = 256,
                shuffle = True,
                validation_data =(X_test, X_test))

#----- VISUALIZATION
visualize(autoencoder, X_test)
```


**shuffle = True:**

The training data is shuffled at the beginning of each epoch to improve generalization and avoid overfitting

## ■ Task 2.2: Autoencoders



- $N\_EPOCHS=2$




- $N\_EPOCHS=20$




## ■ Task 2.3: CNN

- Convolutional neural network




## ■ Task 2.3: CNN



```
#----- RESHAPE
if k.image_data_format() == 'channels_first': # theano
    Xtrain = Xtrain.reshape(Xtrain.shape[0], 1, 28, 28)
    Xtest = Xtest.reshape(Xtest.shape[0], 1, 28, 28)
    inpx = (1, 28, 28)
else:
    Xtrain = Xtrain.reshape(Xtrain.shape[0], 28, 28, 1)
    Xtest = Xtest.reshape(Xtest.shape[0], 28, 28, 1)
    inpx = (28, 28, 1)
```

- Define dimension of the image
  - `k.image_data_format() == 'channels_first'`
  - False in mnist
  - Check how the image data is formatted in the current environment.

## ■ Task 2.3: CNN



```

model = Sequential()

#----- CNN 1
model.add( Conv2D(64, (3, 3), activation='relu') )
model.add( MaxPooling2D(pool_size=(2, 2)))

#----- CNN 2
model.add( Conv2D(32, (3, 3), activation='relu') )
model.add( MaxPooling2D(pool_size=(2, 2)))

#----- VECTOR output
model.add(Flatten())

#----- MLNN
model.add(Dense(14, activation='relu'))
model.add(Dense( 6, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

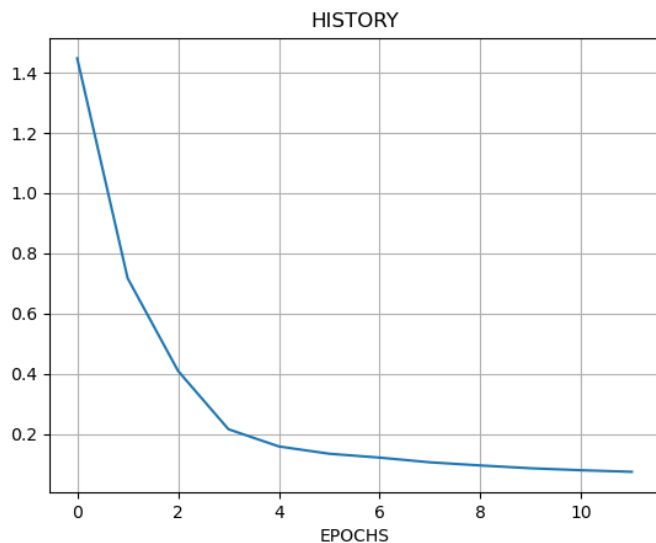

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(Xtrain, Ttrain, epochs=12, batch_size=500)

print(history.history['loss'])

```

## ■ Task 2.3: CNN



EPOCHS=2



EPOCHS=22



```
train = model.predict(Xtrain).round()
Ytest = model.predict(Xtest).round()
scores = model.evaluate(Xtest, Ttest, verbose=0)

print("Baseline Error: %.2f%%" % (100-scores[1]*100))
```



### ■ **Conv2D( 64, (3, 3), activation='relu')**

- Performs 2D convolution,
- Specifies the number of filters (kernels) = 64
- Mask : Indicates the filter size or kernel size, which is 3x3.
- Activation: 'relu'

### ■ **MaxPooling2D(pool\_size=(2, 2))**

- Downsampling the spatial dimensions of feature maps.
- In this case, a 2x2 reducing the spatial dimensions by a factor of 2

### ■ **Flatten()**

- Takes a multi-dimensional tensor (e.g., 2D, 3D, or higher) and flattens it into a single 1D vector.
- Required before feeding the output of a convolutional or pooling layer into a fully connected (dense) layer.

# Contents

- 1 | Objectives
- 2 | Dataset
- 3 | Tasks
- 4 | Conclusions

- Deep learning / CNN
- Concepts
  - 1 | Implement Digit recognition from scratch
  - One layer network ? One hidden layer ‘
- 2 | Use python/keras learning functionalities
  - MLL
  - Autoencoders
  - Convolutional neural network
- 3 | Evaluation the performance of the DNN/CNN classifier